

# T13: CUDA Implementation of Diffusion-Limited Aggregation (DLA) for Procedural Generation of Resources

Ajmal  
cs22b2046@iiitdm.ac.in

March 19, 2025

## 1 Introduction

CUDA implementation of a variation of Diffusion-Limited Aggregation (DLA) simulation for procedural resource generation. The simulation models the behavior of particles (walkers) moving randomly on a grid and aggregating to form clusters, where each cluster indicates a particular type of resource (red and blue in our environment, including purple depending on requirements). The CUDA implementation is compared against a serial implementation. The simulation was executed on an NVIDIA T4 GPU, which has 2560 CUDA cores and 16 GB of memory in Google Colab.

## 2 Performance Comparison

### 2.1 Serial Implementation

- **Execution Time:** 218.856 seconds

### 2.2 CUDA Implementation

- **Execution Time:** 6.179 seconds
- **GPU Used:** NVIDIA T4 (2560 CUDA cores, 16 GB memory)

### 2.3 Speedup

The speedup is calculated as the ratio of the serial execution time (in colab CPU) to the parallel execution time:

$$\text{Speedup} = \frac{\text{Serial Time}}{\text{Parallel Time}} = \frac{218.856}{6.179} \approx 35.419$$

## 2.4 Parallelization Factor

The parallelization factor indicates how effectively the parallel implementation utilizes the available hardware resources. It is calculated as:

For the NVIDIA T4 GPU with 2560 CUDA cores, the parallelization factor is:

$$\text{Parallelization Fraction} \approx 0.97214$$

## 3 Code Explanation

### 3.1 Grid Initialization

The grid is initialized with a fixed width and height, and a seed is placed at the center to start the aggregation process. Modifying the seed placement will give rise to different structures.

```
1 Grid grid = GridInit(width, height);
2 grid.g[(height / 2) * width + (width / 2)] = YELLOW;
3 grid.g[(height / 2) * width + (width / 2) + 1] = YELLOW;
```

### 3.2 Walker Initialization

Walkers are initialized with random positions and colors (in device, using curand for thread safety). The `WalkerInit` function sets up the walker's initial state.

```
1 Walker *h_walkers = (Walker*)malloc(maxWalkers * sizeof(Walker));
2 for (int i = 0; i < maxWalkers; i++) {
3     h_walkers[i] = WalkerInit(0, 0, 0);
4     h_walkers[i].active = 1;
5 }
```

### 3.3 CUDA Kernel for DLA Simulation

The `dlaKernel` function is the core of the CUDA implementation. It handles the random movement of walkers and their interaction with the grid.

```
1 __global__ void dlaKernel(int *grid, int width, int height, Walker
2     *walkers, int numWalkers, int maxSteps, curandState *states) {
3     int idx = blockIdx.x * blockDim.x + threadIdx.x;
4     if (idx >= numWalkers) return;
5
6     curandState localState = states[idx];
7     Walker walker = walkers[idx];
8
9     if (!walker.active) return;
10
11     int color = (idx % 2 == 0) ? RED : BLUE;
12     int edge = curand(&localState) % 4;
13     // Spawn walker at edge
```

```

14     switch (edge) {
15         case 0: walker.x = curand(&localState) % width; walker.y =
0; break;
16         case 1: walker.x = curand(&localState) % width; walker.y =
height - 1; break;
17         case 2: walker.x = 0; walker.y = curand(&localState) %
height; break;
18         case 3: walker.x = width - 1; walker.y = curand(&localState
) % height; break;
19     }
20
21     walker.color = color;
22
23     for (int step = 0; step < maxSteps; step++) {
24         walkerMove(&walker, grid, width, height, &localState);
25         int adjacent = isAdjacentToCluster(grid, width, height, &
walker);
26
27         if (adjacent) {
28             if (isWalkerClrMatch(&walker, YELLOW) ||
isWalkerClrMatch(&walker, PURPLE) || isWalkerClrMatch(&walker,
walker.color)) {
29                 walkerPlaceOnGrid(&walker, grid, width);
30                 spawnGreen(walker, grid, width, height, 1, 0.01, &
localState);
31             } else {
32                 detonate(grid, walker.x, walker.y, width, height,
1, 1.0, 0.4, &localState);
33             }
34             walker.active = 0;
35             break;
36         }
37     }
38
39     walkers[idx] = walker;
40     states[idx] = localState;
41 }

```

### 3.4 Batch Processing

The simulation processes walkers in batches to manage memory and computational load efficiently.

```

1 int batchSize = 1000;
2 int numBatches = (maxWalkers + batchSize - 1) / batchSize;
3
4 for (int batch = 0; batch < numBatches; batch++) {
5     int startIdx = batch * batchSize;
6     int endIdx = min(startIdx + batchSize, maxWalkers);
7     int currentBatchSize = endIdx - startIdx;
8
9     dlaKernel<<<numBlocks, blockSize>>>(d_grid, width, height,
d_walkers + startIdx, currentBatchSize, maxSteps, d_states +
startIdx);
10    cudaDeviceSynchronize();
11 }

```

## 4 Results




Figure 1: CUDA 1




Figure 2: CUDA 2




Figure 3: CUDA 3




Figure 4: CUDA 4




Figure 5: CUDA 5




Figure 6: CUDA 6




Figure 7: CUDA 7




Figure 8: CUDA 8




Figure 9: CUDA 9